

Object Detection and Autopilot Safety analysis

Part 1 : Object Detection

Import libraries

```
import torch
import os
import shutil
import pathlib
from glob import glob
import pandas as pd
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import cv2
import numpy as np
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
```

Define the data paths

```
# For colab execution

zip_file_path = '/content/drive/MyDrive/DL/od.zip'

data_root_colab = '/content/drive/MyDrive/DL/Object Detection'
image_data_path_colab = '/content/drive/MyDrive/DL/Object Detection/Images'
labels_path_colab = f'{data_root_colab}/labels.csv'

#for Kaggle execution
```

```
image_data_path_kaggle = '/kaggle/input/object-detection/Object Detection/Images'  
labels_path_kaggle = '/kaggle/input/object-detection/Object Detection/labels.csv'  
data_root_kaggle = '.'
```

```
image_data_path = None  
data_root = None  
labels_path = None
```

```
if os.path.exists(image_data_path_kaggle):  
    image_data_path = image_data_path_kaggle  
    data_root = data_root_kaggle  
    labels_path = labels_path_kaggle  
    print(f'Using Kaggle data ')  
else:  
    print('Not on kaggle')  
    from google.colab import drive  
    drive.mount('/content/drive')  
    data_root = data_root_colab  
    image_data_path = image_data_path_colab  
    labels_path = labels_path_colab  
    if not os.path.exists(image_data_path):  
        from zipfile import ZipFile  
        zip_ref = ZipFile(zip_file_path, 'r')  
        zip_ref.extractall('/content/drive/MyDrive/DL/')  
        zip_ref.close()  
  
    print(f'Using Google Drive data')  
print(f'image data path is {image_data_path}')  
  
print(f'labels path is at {labels_path}')
```

```
Not on kaggle  
Mounted at /content/drive  
Using Google Drive data  
image data path is /content/drive/MyDrive/DL/Object Detection/Images  
labels path is at /content/drive/MyDrive/DL/Object Detection/labels.csv
```

Create the directory structure for yolo training

```
# create train , validation and test image directories
train_label_path = 'datasets/train/labels'
val_labels_path = 'datasets/val/labels'
train_image_path = 'datasets/train/images'
val_image_path = 'datasets/val/images'
test_image_path = 'datasets/test/images'
```

```
dirs = [train_label_path, val_labels_path, train_image_path, val_image_path, test_image_path]
```

```
for d in dirs:
    if os.path.exists(d):
        shutil.rmtree(d)
    os.makedirs(d)
    print(f'created {d}')

if not os.path.exists('datasets/data.yaml'):
    pathlib.Path.touch('datasets/data.yaml')
```

```
created datasets/train/labels
created datasets/val/labels
created datasets/train/images
created datasets/val/images
created datasets/test/images
```

```
%ls datasets
```

```
data.yaml  test/  train/  val/
```

Load the labels csv file

```
import pandas as pd

label_data = pd.read_csv(labels_path, names=['id', 'name', 'x_min', 'y_min', 'x_max', 'y_max'])
label_data = label_data.drop_duplicates('id')
label_data['id'] = label_data['id'].astype(str).str.zfill(8)
label_data = label_data.iloc[0:len(os.listdir(image_data_path))]
label_data['image'] = sorted(os.listdir(image_data_path))
```

```
print(label_data.shape)
label_data.head()
```

```
(5626, 7)
```

	id	name	x_min	y_min	x_max	y_max	image
0	00000000	pickup_truck	213	34	255	50	00000000.jpg
5	00000001	bus	205	155	568	314	00000001.jpg
11	00000002	motorized_vehicle	334	362	603	479	00000002.jpg
12	00000003	car	480	99	511	120	00000003.jpg
14	00000004	bus	439	32	719	171	00000004.jpg

Estimate Vehicle numbers

```
def copy_image_to_name(row):
    """
    copies an image to the respective folder
    param : row
    """
    name = row['name']
    dest = f'{data_root}/image_root/{name}/'
    if os.path.exists(dest):
        shutil.copy(f'{image_data_path}/{row["image"]}', dest)
    else:
        os.makedirs(dest)
        shutil.copy(f'{image_data_path}/{row["image"]}', dest)
_label_data.apply(copy_image_to_name,axis=1)
```

```
def get_subdirectories(directory) -> list:
    with os.scandir(directory) as entries:
        return [entry.name for entry in entries if entry.is_dir()]

def plot_class_numbers(directory,title) -> pd.DataFrame:
    """
```

```

Plots the number of images in each class
:param directory: the directory containing the images
:param title: the title of the plot
:return: a dataframe containing the number of images in each class
'''

sub_dirs = get_subdirectories(directory)
images = {}
image_class = []
num_images = []
class_dict = {}

for d in sub_dirs:
    image_class.append(d)
    num_images.append(len(os.listdir(str(directory)+'/'+d)))

df = pd.DataFrame({'Image_Class':image_class,'Num_Images':num_images})
df = df.sort_values(by='Num_Images',ascending=True).reset_index()


indices = sorted(df['index'].values.tolist())
class_names = df['Image_Class'].values
for i in indices:
    class_dict[i] = class_names[i]
sns.barplot(data=df,x='Image_Class',y='Num_Images',palette=sns.mpl_palette('Dark2'))
plt.xlabel('vehicle types')
plt.title(f' {title}' )
plt.xticks(rotation = 90)
plt.gca().spines[['top', 'right',]].set_visible(False)
plt.show()
return df

```

```

plot_class_numbers(f'{data_root}/image_root','')

```



	index	Image_Class	Num_Images
0	6	non-motorized_vehicle	32
1	10	motorcycle	38
2	3	bicycle	43
3	5	pedestrian	56
4	9	single_unit_truck	80
5	7	work_van	161
6	1	articulated_truck	195
7	4	bus	244
8	8	motorized_vehicle	383
9	0	pickup_truck	801
10	2	car	3593

There is substantial class imbalance

Create a map of vehicle types to index

```
import pprint
class_names = label_data.name.unique().tolist()
d = {}
for i in range(len(class_names)):
    d[class_names[i]] = i
pprint.pprint(d)
```

```
{'articulated_truck': 4,  
 'bicycle': 8,  
 'bus': 1,  
 'car': 3,  
 'motorcycle': 10,  
 'motorized_vehicle': 2,  
 'non-motorized_vehicle': 7,  
 'pedestrian': 6,  
 'pickup_truck': 0,  
 'single_unit_truck': 9,  
 'work_van': 5}
```

Create labels for training and validation

```
import cv2  
import numpy as np  
import pathlib  
def get_label_filename(row):  
    return row['id']+'.txt'  
  
def normalize_bb(x_min,y_min,x_max,y_max,id):  
    img_path = f'{image_data_path}/{id}.jpg'  
    if os.path.exists(img_path):  
        img = cv2.imread(img_path)  
        image_height,image_width,_ = img.shape  
        x_center = (x_min + x_max)/2  
        y_center = (y_min + y_max)/2  
        width = x_max - x_min  
        height = y_max - y_min  
        normalized_x_center = x_center / image_width  
        normalized_y_center = y_center / image_height  
        normalized_width = width / image_width  
        normalized_height = height / image_height  
        return normalized_x_center, normalized_y_center, normalized_width , normalized_height  
  
def get_label_data(row):  
    dims = normalize_bb(row['x_min'], row['y_min'],row['x_max'],row['y_max'],row['id'])  
    class_id = d.get(row['name'])
```



```

yolo_annotaion = f'{class_id} {dims[0]:.3f} {dims[1]:.3f} {dims[2]:.3f} {dims[0]:.3f}'
return yolo_annotaion

def get_image_width(row):
    img_path = f'{image_data_path}/{row['id']}.jpg'
    img = cv2.imread(img_path)
    return img.shape[1]

def get_image_height(row):
    img_path = f'{image_data_path}/{row['id']}.jpg'
    img = cv2.imread(img_path)
    return img.shape[0]

def create_yaml():
    file_path = 'datasets/data.yaml'
    if os.path.exists(file_path):
        os.remove(file_path)

    train = 'train: train'
    val = 'val: val'
    nc = "nc: 11"
    names = f"names: {str(class_names)}"
    lines_to_append = [train, val, nc, names]

    try:
        with open(file_path, 'a') as file:
            for line in lines_to_append:
                file.write(f'{line}\n')
            print('data.yaml created')
    except Exception as e:
        print(f"An error occurred: {e}")

```

Create data.yaml for training

```
create_yaml()
```

```
data.yaml created
```

```
%cat datasets/data.yaml
```

```
train: train
val: val
nc: 11
names: ['pickup_truck', 'bus', 'motorized_vehicle', 'car', 'articulated_truck', 'work_van', 'pedestrian', 'non-motorized_vehicle', 'bicycle', 'single_unit_truck', 'motorcycle']
```

Create annotations in YOLO5 format

```
from glob import glob
import warnings
warnings.filterwarnings('ignore')

num_images = len(glob(f'{image_data_path}/*.jpg'))
num_train_images = int(num_images*0.8)

num_test_images = 60
num_val_images = num_images - num_train_images - num_test_images

print(f'total :{num_images}, train : {num_train_images}, val :{num_val_images}, test:{num_test_images}')
```

labels_train = label_data.iloc[0:num_train_images]
labels_val = label_data.iloc[num_train_images:num_images-num_test_images]
print(f'train shape : {labels_train.shape}, val shape: {labels_val.shape}')

total :5626, train : 4500, val :1066, test:60
train shape : (4500, 7), val shape: (1066, 7)

```
labels_train['filename'] = labels_train.apply(get_label_filename,axis=1)
labels_train['text'] = labels_train.apply(get_label_data,axis=1)
labels_train['image_width'] = labels_train.apply( get_image_width,axis=1)
labels_train['image_height'] = labels_train.apply( get_image_height,axis=1)

labels_val['filename'] = labels_val.apply( get_label_filename,axis=1)
labels_val['text'] = labels_val.apply( get_label_data,axis=1)
labels_val['image_width'] = labels_val.apply( get_image_width,axis=1)
labels_val['image_height'] = labels_val.apply( get_image_height,axis=1)
```

```
print(f'train shape : {labels_train.shape}, val shape: {labels_val.shape}')
labels_train.sample(10)
```

```
train shape : (4500, 11), val shape: (1066, 11)
```

	id	name	x_min	y_min	x_max	y_max	image	filename	text	image_width	image_height
13607	00004345	car	135	259	192	348	00004345.jpg	00004345.txt	3 0.227 0.632 0.079 0.227	720	480
3157	00001001	car	607	147	654	172	00001001.jpg	00001001.txt	3 0.876 0.332 0.065 0.876	720	480
9668	00003069	car	129	62	203	97	00003069.jpg	00003069.txt	3 0.485 0.349 0.216 0.485	342	228
8476	00002708	car	120	94	198	144	00002708.jpg	00002708.txt	3 0.465 0.522 0.228 0.465	342	228
5320	00001672	car	150	172	225	219	00001672.jpg	00001672.txt	3 0.260 0.407 0.104 0.260	720	480
5781	00001828	bicycle	357	136	394	236	00001828.jpg	00001828.txt	8 0.522 0.388 0.051 0.522	720	480
684	00000211	motorized_vehicle	0	244	72	351	00000211.jpg	00000211.txt	2 0.050 0.620 0.100 0.050	720	480
7226	00002261	car	514	159	715	292	00002261.jpg	00002261.txt	3 0.853 0.470 0.279 0.853	720	480
6702	00002114	car	70	130	219	201	00002114.jpg	00002114.txt	3 0.201 0.345 0.207 0.201	720	480
4456	00001412	pickup_truck	42	47	144	91	00001412.jpg	00001412.txt	0 0.272 0.303 0.298 0.272	342	228

Define function to create the yolo label file

```

from pathlib import Path

# create the label file
def create_label_file(label_path,file_name,text):

    f = Path(f'{label_path}/{file_name}')
    f.write_text(text)

```

Create label files for train and validation images

```

labels_train.apply(lambda row: create_label_file(train_label_path,row['filename'],row['text']) , axis =1)
_labels_val.apply(lambda row: create_label_file(val_labels_path,row['filename'],row['text']) , axis =1)

```

Copy the images to train, val and test folders

```

# copy train images
from tqdm import tqdm

def _copy(src,dest,file_name):
    if not src.startswith('.'):
        src_file = f'{src}/{file_name}'
        if not os.path.isdir(src_file):
            shutil.copy(src_file,dest)

num_train_images = labels_train.shape[0]
num_val_images = labels_val.shape[0]

images = sorted(os.listdir(image_data_path))

print(num_train_images,num_train_images+num_val_images,num_test_images)
print('Copying training images')
for i in tqdm(range(0,num_train_images)):
    _copy(image_data_path,f'{train_image_path}/{',images[i])

print('Copying validation images')

```

```

for i in tqdm(range(num_train_images,num_train_images+num_val_images)):
    _copy(image_data_path,f'{val_image_path}/',images[i])

print('copying test imges')
for i in tqdm(range(num_train_images+num_val_images,num_images)):
    _copy(image_data_path,f'{test_image_path}/',images[i])

```

4500 5566 60

Copying training images

100%|██████████| 4500/4500 [00:08<00:00, 504.77it/s]

Copying validation images

100%|██████████| 1066/1066 [00:01<00:00, 533.72it/s]

copying test imges

100%|██████████| 60/60 [00:00<00:00, 542.26it/s]

install ultralytics

```
!pip install ultralytics > install.txt
```

Train the model

```

from ultralytics import YOLO
data_yaml_kaggle = "/kaggle/working/datasets/data.yaml"
data_yaml_colab = "/content/datasets/data.yaml"
epochs = 100
patience = 10
data_yaml = data_yaml_kaggle if os.path.exists(data_yaml_kaggle) else data_yaml_colab
model = YOLO("yolov5su.pt")
train_results=model.train(data=data_yaml, epochs=epochs, patience = patience, imgsz=640)

```

Ultralytics 8.3.107 🚀 Python-3.11.11 torch-2.5.1+cu124 CUDA:0 (Tesla T4, 15095MiB)

engine/trainer: task=detect, mode=train, model=yolov5su.pt, data=/kaggle/working/datasets/data.yaml, epochs=100, time=None, patience=10, batch=16, imgsz=640, save=True, save_period=-1, cache=False, device=None, workers=8, project=None, name=train3, exist_ok=False, pretrained=True, optimizer=auto, verbose=True, seed=0, deterministic=True, single_cls=False, rect=False, cos_lr=False, close_mosaic=10, resume=False, amp=True, fraction=1.0, profile=False, freeze=None, multi_scale=False, overlap_mask=True, mask_ratio=4, dropout=0.0, val=True, split=val, save_json=False, conf=None, iou=0.7, max_det=300, half=False, dnn=False, plots=True, source=None, vid_stride=1, stream_buffer=False, visualize=False, augment=False, agnostic_nms=False, classes=None, retina_masks=False, embed=None, show=False, save_frames=False, save_txt=False, save_conf=False, save_crop=False, show_labels=True, show_conf=True, show_boxes=True, line_width=None, format=torchscript, keras=False, optimize=False, int8=False, dynamic=False, simplify=True, opset=None, workspace=None, nms=False, lr0=0.01, lrf=0.01, momentum=0.937, weight_decay=0.0005, warmup_epochs=3.0, warmup_momentum=0.8, warmup_bias_lr=0.1, box=7.5, cls=0.5, dfl=1.5, pose=12.0, kobj=1.0, nbs=64, hsv_h=0.015, hsv_s=0.7, hsv_v=0.4, degrees=0.0, translate=0.1, scale=0.5, shear=0.0, perspective=0.0, flipud=0.0, fliplr=0.5, bgr=0.0, mosaic=1.0, mixup=0.0, copy_paste=0.0, copy_paste_mode=flip, auto_augment=randaugment, erasing=0.4, crop_fraction=1.0, cfg=None, tracker=botsort.yaml, save_dir=runs/detect/train3

Overriding model.yaml nc=80 with nc=11

	from	n	params	module	arguments
0		-1 1	3520	ultralytics.nn.modules.conv.Conv	[3, 32, 6, 2, 2]
1		-1 1	18560	ultralytics.nn.modules.conv.Conv	[32, 64, 3, 2]
2		-1 1	18816	ultralytics.nn.modules.block.C3	[64, 64, 1]
3		-1 1	73984	ultralytics.nn.modules.conv.Conv	[64, 128, 3, 2]
4		-1 2	115712	ultralytics.nn.modules.block.C3	[128, 128, 2]
5		-1 1	295424	ultralytics.nn.modules.conv.Conv	[128, 256, 3, 2]
6		-1 3	625152	ultralytics.nn.modules.block.C3	[256, 256, 3]
7		-1 1	1180672	ultralytics.nn.modules.conv.Conv	[256, 512, 3, 2]
8		-1 1	1182720	ultralytics.nn.modules.block.C3	[512, 512, 1]
9		-1 1	656896	ultralytics.nn.modules.block.SPPF	[512, 512, 5]
10		-1 1	131584	ultralytics.nn.modules.conv.Conv	[512, 256, 1, 1]
11		-1 1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
12	[-1, 6]	1	0	ultralytics.nn.modules.conv.Concat	[1]
13		-1 1	361984	ultralytics.nn.modules.block.C3	[512, 256, 1, False]
14		-1 1	33024	ultralytics.nn.modules.conv.Conv	[256, 128, 1, 1]
15		-1 1	0	torch.nn.modules.upsampling.Upsample	[None, 2, 'nearest']
16	[-1, 4]	1	0	ultralytics.nn.modules.conv.Concat	[1]
17		-1 1	90880	ultralytics.nn.modules.block.C3	[256, 128, 1, False]
18		-1 1	147712	ultralytics.nn.modules.conv.Conv	[128, 128, 3, 2]
19	[-1, 14]	1	0	ultralytics.nn.modules.conv.Concat	[1]
20		-1 1	296448	ultralytics.nn.modules.block.C3	[256, 256, 1, False]
21		-1 1	590336	ultralytics.nn.modules.conv.Conv	[256, 256, 3, 2]
22	[-1, 10]	1	0	ultralytics.nn.modules.conv.Concat	[1]
23		-1 1	1182720	ultralytics.nn.modules.block.C3	[512, 512, 1, False]
24	[17, 20, 23]	1	2120305	ultralytics.nn.modules.head.Detect	[11, [128, 256, 512]]

YOLOv5s summary: 153 layers, 9,126,449 parameters, 9,126,433 gradients, 24.1 GFLOPs

Transferred 421/427 items from pretrained weights

TensorBoard: Start with 'tensorboard --logdir runs/detect/train3', view at http://localhost:6006/

Freezing layer 'model.24.dfl.conv.weight'

AMP: running Automatic Mixed Precision (AMP) checks...

AMP: checks passed ✓

train: Scanning /kaggle/working/datasets/train/labels.cache... 4500 images, 0 backgrounds, 0 corrupt: 100%|██████████| 4500/4500 [00:00<?, ?it/s]

albumentations: Blur(p=0.01, blur_limit=(3, 7)), MedianBlur(p=0.01, blur_limit=(3, 7)), ToGray(p=0.01, num_output_channels=3, method='weighted_average'), CLAHE(p=0.01, clip_limit=(1.0, 4.0), tile_grid_size=(8, 8))

val: Scanning /kaggle/working/datasets/val/labels.cache... 1066 images, 0 backgrounds, 0 corrupt: 100%|██████████| 1066/1066 [00:00<?, ?it/s]

Plotting labels to runs/detect/train3/labels.jpg...

optimizer: 'optimizer=auto' found, ignoring 'lr0=0.01' and 'momentum=0.937' and determining best 'optimizer', 'lr0' and 'momentum' automatically...

optimizer: AdamW(lr=0.000667, momentum=0.9) with parameter groups 69 weight(decay=0.0), 76 weight(decay=0.0005), 75 bias(decay=0.0)

TensorBoard: model graph visualization added ✓

Image sizes 640 train, 640 val

Using 2 dataloader workers

Logging results to runs/detect/train3

Starting training for 100 epochs...

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
1/100	6.65G	1.887	4.267	1.829	12	640: 100% ██████████ 282/282 [01:13<00:00, 3.83it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.28it/s]
	all	1066	1066	0.683	0.14	0.0755 0.0385

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
2/100	6.67G	1.631	2.714	1.666	11	640: 100% ██████████ 282/282 [01:11<00:00, 3.97it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.30it/s]
	all	1066	1066	0.615	0.0942	0.0654 0.0354

Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
3/100	6.67G	1.605	2.582	1.647	11	640: 100% ██████████ 282/282 [01:11<00:00, 3.97it/s]
	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.19it/s]

	all	1066	1066	0.606	0.0992	0.0592	0.0328	
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
s]	4/100	6.67G	1.548	2.476	1.614	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.22it/s]	
	all	1066	1066	0.522	0.136	0.0906	0.051	
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
s]	5/100	6.67G	1.481	2.368	1.554	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.12it/s]	
	all	1066	1066	0.716	0.112	0.0891	0.0574	
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
s]	6/100	6.67G	1.422	2.242	1.52	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.24it/s]	
	all	1066	1066	0.733	0.114	0.124	0.0804	
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
s]	7/100	6.67G	1.376	2.219	1.488	8	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.24it/s]	
	all	1066	1066	0.624	0.126	0.107	0.07	
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
s]	8/100	6.67G	1.338	2.132	1.459	5	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.18it/s]	
	all	1066	1066	0.69	0.196	0.156	0.0965	
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
s]	9/100	6.67G	1.293	2.081	1.428	11	640: 100% ██████████ 282/282 [01:10<00:00, 4.01it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.29it/s]	
	all	1066	1066	0.395	0.187	0.14	0.092	
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size	
s]	10/100	6.67G	1.265	2.039	1.412	8	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.18it/s]	
	all	1066	1066	0.461	0.248	0.174	0.119	

	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	11/100	6.67G	1.247	2.026	1.392	12	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class		Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.24it/s]
	all		1066	1066	0.725	0.149	0.181 0.126
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	12/100	6.67G	1.219	1.952	1.373	12	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]
	Class		Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.22it/s]
	all		1066	1066	0.412	0.257	0.171 0.121
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	13/100	6.67G	1.219	1.929	1.364	7	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class		Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.25it/s]
	all		1066	1066	0.499	0.287	0.197 0.133
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	14/100	6.67G	1.169	1.855	1.33	6	640: 100% ██████████ 282/282 [01:10<00:00, 4.00it/s]
	Class		Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.28it/s]
	all		1066	1066	0.405	0.216	0.196 0.139
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	15/100	6.67G	1.175	1.862	1.333	8	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class		Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.19it/s]
	all		1066	1066	0.347	0.304	0.217 0.16
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	16/100	6.67G	1.145	1.827	1.311	5	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class		Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.27it/s]
	all		1066	1066	0.386	0.331	0.235 0.154
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	17/100	6.67G	1.132	1.801	1.305	10	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]
	Class		Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.21it/s]
	all		1066	1066	0.33	0.35	0.243 0.169
	Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size

s]	18/100	6.67G	1.124	1.775	1.294	7	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.32it/s]
	all		1066	1066	0.421	0.282	0.238	0.16
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	19/100	6.67G	1.113	1.762	1.285	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.30it/s]
	all		1066	1066	0.288	0.337	0.246	0.175
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	20/100	6.67G	1.089	1.743	1.27	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.25it/s]
	all		1066	1066	0.49	0.339	0.24	0.168
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	21/100	6.67G	1.07	1.725	1.256	11	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.41it/s]
	all		1066	1066	0.575	0.243	0.248	0.168
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	22/100	6.67G	1.085	1.695	1.268	12	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.26it/s]
	all		1066	1066	0.443	0.278	0.268	0.189
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	23/100	6.67G	1.041	1.673	1.244	4	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.40it/s]
	all		1066	1066	0.321	0.381	0.258	0.184
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	24/100	6.67G	1.058	1.684	1.248	10	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.33it/s]
	all		1066	1066	0.506	0.339	0.292	0.208
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	

s]	25/100	6.67G	1.043	1.65	1.239	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.00it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.27it/s]
	all	1066	1066	0.243	0.376	0.267	0.2
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	26/100	6.67G	1.004	1.594	1.215	9	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.44it/s]
	all	1066	1066	0.338	0.375	0.302	0.213
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	27/100	6.67G	1.024	1.646	1.223	7	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.31it/s]
	all	1066	1066	0.386	0.281	0.289	0.199
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	28/100	6.67G	0.9985	1.607	1.21	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.01it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.35it/s]
	all	1066	1066	0.466	0.349	0.293	0.211
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	29/100	6.67G	0.9956	1.588	1.206	11	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.20it/s]
	all	1066	1066	0.425	0.374	0.296	0.214
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	30/100	6.67G	0.994	1.59	1.206	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.14it/s]
	all	1066	1066	0.341	0.389	0.299	0.207
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	31/100	6.67G	0.9869	1.575	1.195	8	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.42it/s]
	all	1066	1066	0.419	0.41	0.278	0.216
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size

s]	32/100	6.67G	0.9588	1.529	1.179	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.30it/s]
	all	1066	1066	0.242	0.39	0.28	0.216
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	33/100	6.67G	0.9705	1.541	1.185	5	640: 100% ██████████ 282/282 [01:10<00:00, 4.01it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.24it/s]
	all	1066	1066	0.396	0.389	0.322	0.245
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	34/100	6.67G	0.958	1.517	1.183	6	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.36it/s]
	all	1066	1066	0.347	0.383	0.311	0.244
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	35/100	6.67G	0.9581	1.541	1.184	10	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.21it/s]
	all	1066	1066	0.331	0.41	0.29	0.221
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	36/100	6.67G	0.9421	1.508	1.171	5	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.34it/s]
	all	1066	1066	0.26	0.417	0.3	0.234
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	37/100	6.67G	0.9267	1.483	1.158	7	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.32it/s]
	all	1066	1066	0.478	0.356	0.308	0.24
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	38/100	6.67G	0.9275	1.468	1.157	10	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.19it/s]
	all	1066	1066	0.468	0.346	0.321	0.252
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size

s]	39/100	6.67G	0.929	1.482	1.156	7	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.36it/s]
	all	1066	1066	0.434	0.39	0.326	0.241	
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	40/100	6.67G	0.9107	1.433	1.147	5	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.20it/s]
	all	1066	1066	0.352	0.418	0.311	0.238	
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	41/100	6.67G	0.9006	1.426	1.141	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.35it/s]
	all	1066	1066	0.285	0.382	0.321	0.254	
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	42/100	6.67G	0.9028	1.436	1.14	8	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.44it/s]
	all	1066	1066	0.372	0.356	0.334	0.262	
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	43/100	6.67G	0.8907	1.418	1.135	7	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.29it/s]
	all	1066	1066	0.411	0.459	0.341	0.265	
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	44/100	6.67G	0.8859	1.416	1.128	5	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.43it/s]
	all	1066	1066	0.347	0.452	0.354	0.275	
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	
s]	45/100	6.67G	0.8805	1.381	1.13	7	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]	
	Class		Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.25it/s]
	all	1066	1066	0.288	0.447	0.331	0.257	
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size	

s]	46/100	6.67G	0.8838	1.387	1.126	6	640: 100% ██████████ 282/282 [01:10<00:00, 4.01it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.32it/s]
	all	1066	1066	0.448	0.343	0.305	0.236
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	47/100	6.67G	0.8759	1.373	1.119	5	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.49it/s]
	all	1066	1066	0.393	0.357	0.329	0.257
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	48/100	6.67G	0.8776	1.393	1.124	6	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.47it/s]
	all	1066	1066	0.274	0.456	0.323	0.254
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	49/100	6.67G	0.8612	1.352	1.115	11	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.31it/s]
	all	1066	1066	0.403	0.451	0.36	0.281
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	50/100	6.67G	0.8492	1.337	1.105	8	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.28it/s]
	all	1066	1066	0.325	0.422	0.353	0.271
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	51/100	6.67G	0.8587	1.326	1.116	9	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.36it/s]
	all	1066	1066	0.452	0.379	0.353	0.283
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
s]	52/100	6.67G	0.8389	1.334	1.099	9	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.43it/s]
	all	1066	1066	0.435	0.39	0.339	0.269
	Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size

s]	53/100	6.67G	0.8477	1.309	1.107	4	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.46it/s]	
		all	1066	1066	0.365	0.411	0.354	0.276
		Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	54/100	6.67G	0.8392	1.309	1.107	4	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.40it/s]	
		all	1066	1066	0.364	0.4	0.354	0.278
		Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	55/100	6.67G	0.8323	1.305	1.095	5	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.27it/s]	
		all	1066	1066	0.276	0.427	0.327	0.253
		Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	56/100	6.67G	0.8401	1.293	1.095	6	640: 100% ██████████ 282/282 [01:09<00:00, 4.03it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.09it/s]	
		all	1066	1066	0.469	0.346	0.337	0.269
		Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	57/100	6.67G	0.8336	1.29	1.089	6	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.33it/s]	
		all	1066	1066	0.288	0.417	0.348	0.275
		Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	58/100	6.67G	0.8166	1.266	1.087	11	640: 100% ██████████ 282/282 [01:10<00:00, 4.03it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.44it/s]	
		all	1066	1066	0.341	0.431	0.349	0.271
		Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size
s]	59/100	6.67G	0.8017	1.245	1.073	9	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]	
	Class	Images	Instances	Box(P	R	mAP50	mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.35it/s]	
		all	1066	1066	0.477	0.39	0.353	0.268
		Epoch	GPU_mem	box_loss	cls_loss	dfl_loss	Instances	Size

60/100	6.67G	0.808	1.234	1.079	7	640: 100% ██████████ 282/282 [01:10<00:00, 4.02it/s]
s]	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.23it/s]
	all	1066	1066	0.307	0.404	0.336 0.264
Epoch	GPU_mem	box_loss	cls_loss	df1_loss	Instances	Size
61/100	6.67G	0.7981	1.24	1.076	4	640: 100% ██████████ 282/282 [01:10<00:00, 4.00it/s]
s]	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 5.28it/s]
	all	1066	1066	0.498	0.362	0.347 0.268

EarlyStopping: Training stopped early as no improvement observed in last 10 epochs. Best results observed at epoch 51, best model saved as best.pt.

To update EarlyStopping(patience=10) pass a new patience value, i.e. `patience=300` or use `patience=0` to disable EarlyStopping.

61 epochs completed in 1.307 hours.

Optimizer stripped from runs/detect/train3/weights/last.pt, 18.5MB

Optimizer stripped from runs/detect/train3/weights/best.pt, 18.5MB

Validating runs/detect/train3/weights/best.pt...

Ultralytics 8.3.107 🚀 Python-3.11.11 torch-2.5.1+cu124 CUDA:0 (Tesla T4, 15095MiB)

YOLOv5s summary (fused): 84 layers, 9,115,793 parameters, 0 gradients, 23.8 GFLOPs

s]	Class	Images	Instances	Box(P	R	mAP50 mAP50-95): 100% ██████████ 34/34 [00:06<00:00, 4.93it/s]
	all	1066	1066	0.452	0.378	0.353 0.283
	pickup_truck	154	154	0.436	0.558	0.488 0.424
	bus	51	51	0.573	0.784	0.63 0.519
	motorized_vehicle	74	74	0.181	0.0811	0.0714 0.0482
	car	667	667	0.444	0.594	0.518 0.428
	articulated_truck	28	28	0.367	0.56	0.465 0.37
	work_van	39	39	0.491	0.308	0.359 0.305
	pedestrian	13	13	0.293	0.154	0.136 0.0641
	non-motorized_vehicle	6	6	1	0	0.00137 0.000683
	bicycle	8	8	0.418	0.272	0.382 0.294
	single_unit_truck	18	18	0.356	0.222	0.173 0.134
	motorcycle	8	8	0.419	0.625	0.663 0.527

Speed: 0.1ms preprocess, 2.6ms inference, 0.0ms loss, 1.0ms postprocess per image

Results saved to runs/detect/train3

```
from glob import glob
def get_res_dir():
    res_dir_count = len(glob('runs/detect/*'))
```

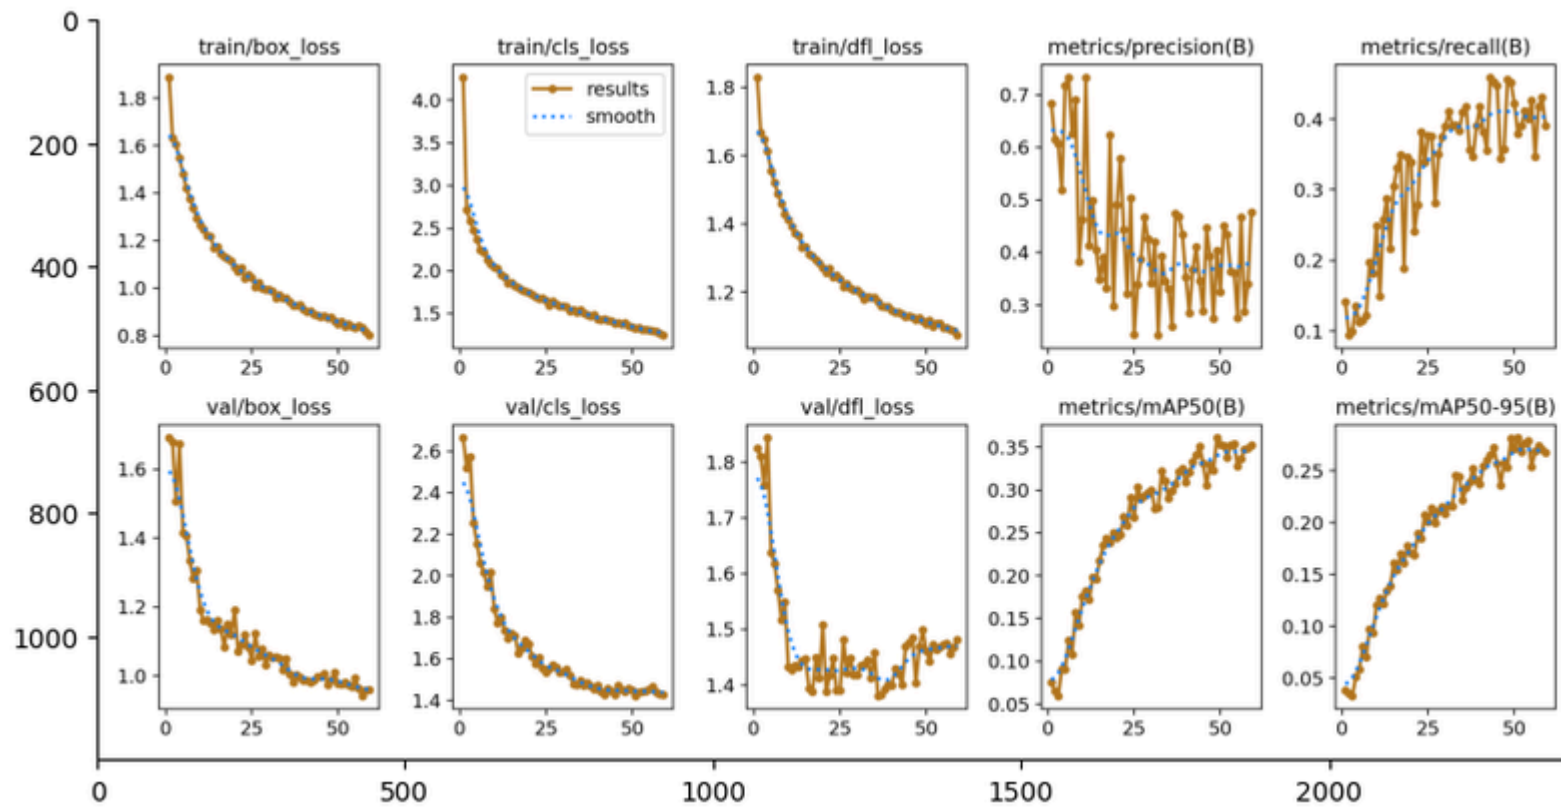


```
# print(f'current number of result directories : {res_dir_count}')
res_dir = f'runs/detect/train{res_dir_count}'
# print(f"result dir : {res_dir}")
return res_dir
```

Precision,Recall,Accuracy curves

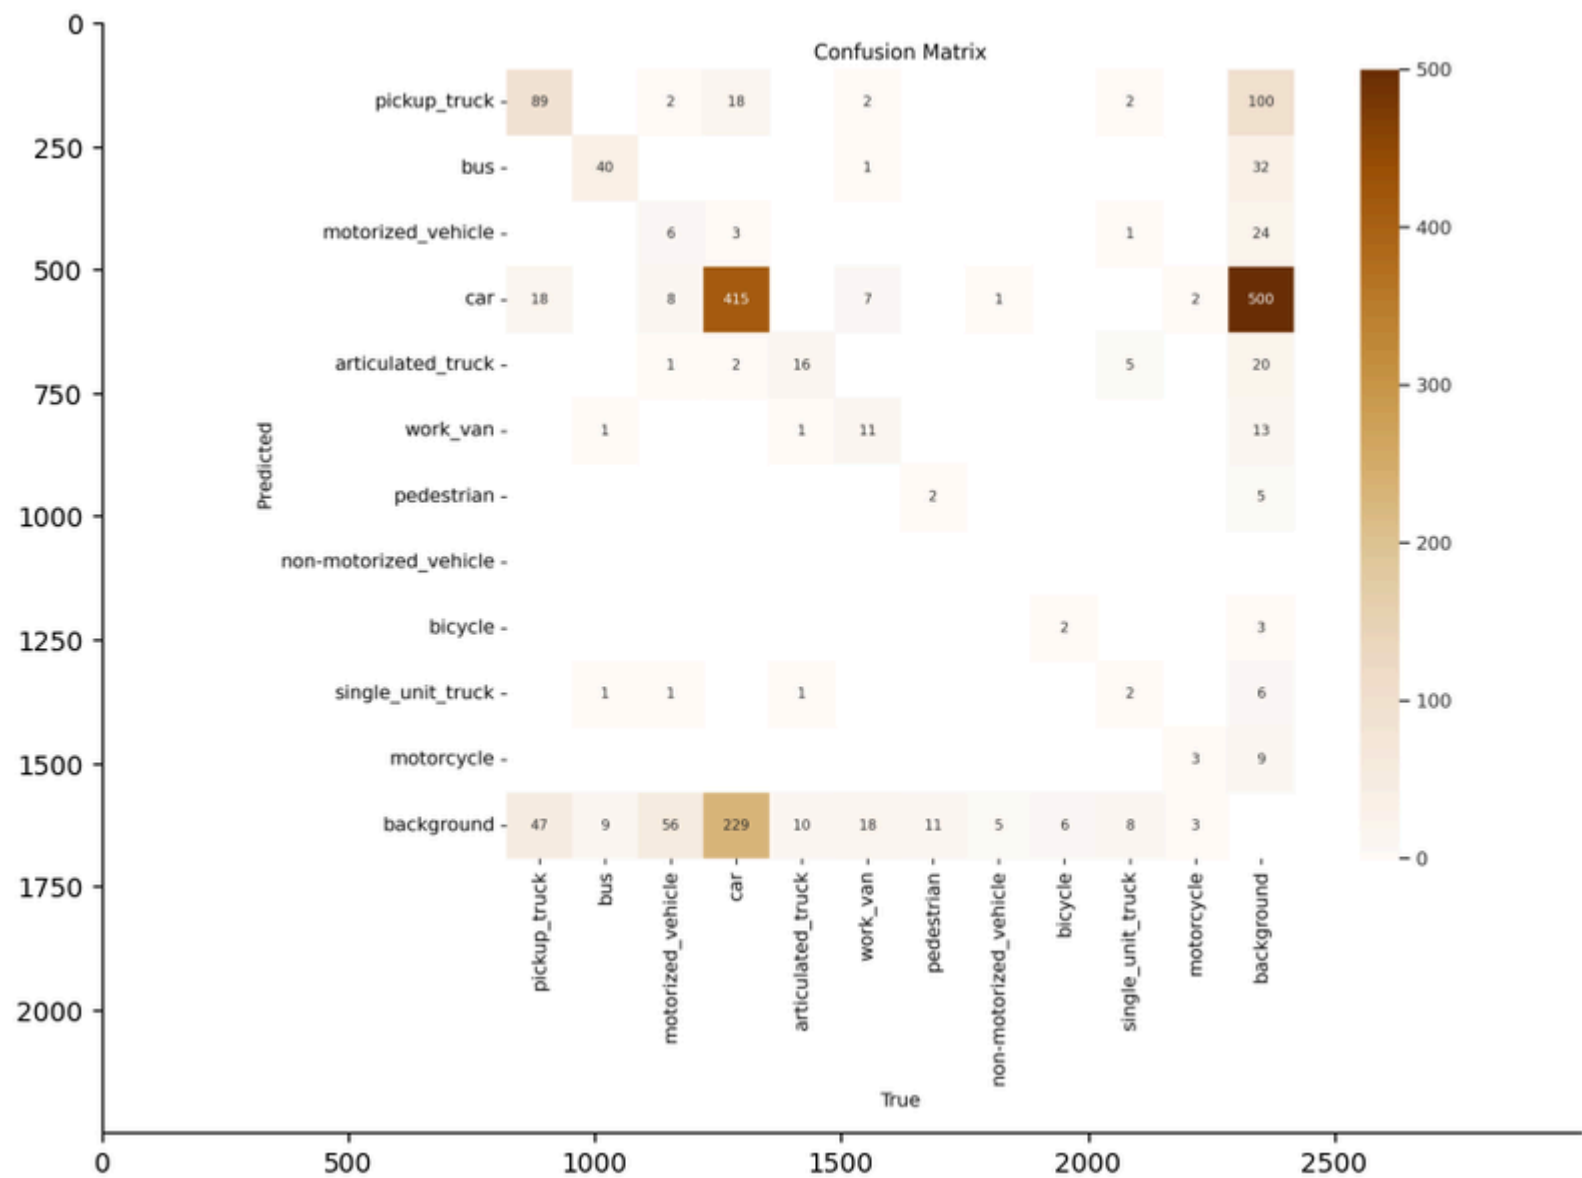
```
import cv2

plt.figure(figsize = (10,10))
res_dir = get_res_dir()
img = cv2.imread(f'{res_dir}/results.png')
plt.gca().spines[['top', 'right',,]].set_visible(False)
_=plt.imshow(img)
```




Confusion Matrix

```
plt.figure(figsize = (10,10))
img = cv2.imread(f'{res_dir}/confusion_matrix.png')
plt.gca().spines[['top', 'right',]].set_visible(False)
_=plt.imshow(img)
```



Confusion Matrix Normalized

```
plt.figure(figsize = (10,10))
img = cv2.imread(f'{res_dir}/confusion_matrix_normalized.png')
plt.gca().spines[['top', 'right',]].set_visible(False)
_=plt.imshow(img)
```



Derive predictions

```
results = model.predict(test_image_path)
```

image 1/60 /kaggle/working/datasets/test/images/00005598.jpg: 448x640 1 work_van, 11.9ms
image 2/60 /kaggle/working/datasets/test/images/00005599.jpg: 448x640 1 bus, 11.8ms
image 3/60 /kaggle/working/datasets/test/images/00005600.jpg: 448x640 1 car, 11.7ms
image 4/60 /kaggle/working/datasets/test/images/00005601.jpg: 448x640 (no detections), 11.7ms
image 5/60 /kaggle/working/datasets/test/images/00005602.jpg: 448x640 1 car, 11.8ms
image 6/60 /kaggle/working/datasets/test/images/00005603.jpg: 448x640 (no detections), 11.8ms
image 7/60 /kaggle/working/datasets/test/images/00005604.jpg: 448x640 1 pickup_truck, 1 car, 11.7ms
image 8/60 /kaggle/working/datasets/test/images/00005605.jpg: 448x640 1 pickup_truck, 11.8ms
image 9/60 /kaggle/working/datasets/test/images/00005606.jpg: 448x640 1 car, 11.8ms
image 10/60 /kaggle/working/datasets/test/images/00005607.jpg: 448x640 2 cars, 9.2ms
image 11/60 /kaggle/working/datasets/test/images/00005608.jpg: 448x640 (no detections), 9.2ms
image 12/60 /kaggle/working/datasets/test/images/00005609.jpg: 448x640 1 bus, 9.2ms
image 13/60 /kaggle/working/datasets/test/images/00005610.jpg: 448x640 1 car, 9.3ms
image 14/60 /kaggle/working/datasets/test/images/00005611.jpg: 448x640 1 pickup_truck, 3 cars, 1 work_van, 9.2ms
image 15/60 /kaggle/working/datasets/test/images/00005612.jpg: 448x640 2 pickup_trucks, 9.2ms
image 16/60 /kaggle/working/datasets/test/images/00005613.jpg: 448x640 2 cars, 9.2ms
image 17/60 /kaggle/working/datasets/test/images/00005614.jpg: 448x640 (no detections), 9.2ms
image 18/60 /kaggle/working/datasets/test/images/00005615.jpg: 448x640 1 motorized_vehicle, 2 cars, 9.2ms
image 19/60 /kaggle/working/datasets/test/images/00005616.jpg: 448x640 1 car, 9.2ms
image 20/60 /kaggle/working/datasets/test/images/00005617.jpg: 448x640 1 pickup_truck, 1 car, 9.2ms
image 21/60 /kaggle/working/datasets/test/images/00005618.jpg: 448x640 (no detections), 9.3ms
image 22/60 /kaggle/working/datasets/test/images/00005619.jpg: 448x640 1 car, 9.2ms
image 23/60 /kaggle/working/datasets/test/images/00005620.jpg: 448x640 1 car, 9.2ms
image 24/60 /kaggle/working/datasets/test/images/00005621.jpg: 448x640 (no detections), 6.6ms
image 25/60 /kaggle/working/datasets/test/images/00005622.jpg: 448x640 2 cars, 6.6ms
image 26/60 /kaggle/working/datasets/test/images/00005623.jpg: 448x640 1 car, 6.5ms
image 27/60 /kaggle/working/datasets/test/images/00005624.jpg: 448x640 2 cars, 6.4ms
image 28/60 /kaggle/working/datasets/test/images/00005625.jpg: 448x640 1 car, 6.4ms
image 29/60 /kaggle/working/datasets/test/images/00005626.jpg: 448x640 1 bus, 6.3ms
image 30/60 /kaggle/working/datasets/test/images/00005627.jpg: 448x640 (no detections), 6.2ms
image 31/60 /kaggle/working/datasets/test/images/00005628.jpg: 448x640 2 cars, 1 pedestrian, 6.4ms
image 32/60 /kaggle/working/datasets/test/images/00005629.jpg: 448x640 3 cars, 6.3ms
image 33/60 /kaggle/working/datasets/test/images/00005630.jpg: 448x640 1 pickup_truck, 6.4ms
image 34/60 /kaggle/working/datasets/test/images/00005631.jpg: 448x640 1 pickup_truck, 1 car, 7.0ms
image 35/60 /kaggle/working/datasets/test/images/00005632.jpg: 448x640 (no detections), 6.5ms
image 36/60 /kaggle/working/datasets/test/images/00005633.jpg: 448x640 1 car, 6.9ms
image 37/60 /kaggle/working/datasets/test/images/00005634.jpg: 448x640 1 car, 6.6ms
image 38/60 /kaggle/working/datasets/test/images/00005635.jpg: 448x640 1 car, 6.5ms
image 39/60 /kaggle/working/datasets/test/images/00005636.jpg: 448x640 1 bus, 2 cars, 6.5ms
image 40/60 /kaggle/working/datasets/test/images/00005637.jpg: 448x640 (no detections), 6.7ms
image 41/60 /kaggle/working/datasets/test/images/00005638.jpg: 448x640 2 articulated_trucks, 6.7ms

```

image 42/60 /kaggle/working/datasets/test/images/00005639.jpg: 448x640 (no detections), 6.8ms
image 43/60 /kaggle/working/datasets/test/images/00005640.jpg: 448x640 4 cars, 6.7ms
image 44/60 /kaggle/working/datasets/test/images/00005641.jpg: 448x640 1 bus, 6.5ms
image 45/60 /kaggle/working/datasets/test/images/00005642.jpg: 448x640 2 cars, 6.4ms
image 46/60 /kaggle/working/datasets/test/images/00005643.jpg: 448x640 2 cars, 6.7ms
image 47/60 /kaggle/working/datasets/test/images/00005644.jpg: 448x640 2 cars, 6.5ms
image 48/60 /kaggle/working/datasets/test/images/00005645.jpg: 448x640 1 car, 6.4ms
image 49/60 /kaggle/working/datasets/test/images/00005646.jpg: 448x640 1 pickup_truck, 2 cars, 6.5ms
image 50/60 /kaggle/working/datasets/test/images/00005647.jpg: 448x640 1 pickup_truck, 6.5ms
image 51/60 /kaggle/working/datasets/test/images/00005648.jpg: 448x640 1 car, 6.7ms
image 52/60 /kaggle/working/datasets/test/images/00005649.jpg: 448x640 1 pickup_truck, 1 bus, 6.3ms
image 53/60 /kaggle/working/datasets/test/images/00005650.jpg: 448x640 1 car, 6.6ms
image 54/60 /kaggle/working/datasets/test/images/00005651.jpg: 448x640 (no detections), 6.4ms
image 55/60 /kaggle/working/datasets/test/images/00005652.jpg: 448x640 (no detections), 6.7ms
image 56/60 /kaggle/working/datasets/test/images/00005653.jpg: 448x640 2 cars, 6.5ms
image 57/60 /kaggle/working/datasets/test/images/00005654.jpg: 448x640 (no detections), 6.5ms
image 58/60 /kaggle/working/datasets/test/images/00005655.jpg: 448x640 1 car, 6.8ms
image 59/60 /kaggle/working/datasets/test/images/00005656.jpg: 448x640 2 cars, 6.3ms
image 60/60 /kaggle/working/datasets/test/images/00005657.jpg: 448x640 2 cars, 6.5ms
Speed: 1.7ms preprocess, 8.0ms inference, 0.9ms postprocess per image at shape (1, 3, 448, 640)

```

```

def plot_box(image,boxes,labels):
    h,w,_ = image.shape
    for box_num,box in enumerate(boxes):
        xyxy = box.xyxy.int().tolist()
        xmin,ymin,xmax,ymax = xyxy[0][0], xyxy[0][1],xyxy[0][2],xyxy[0][3]
        width = xmax-xmin
        height = ymax-ymin
        class_name = class_names[int(labels[box_num])]
        color = COLORS[int(labels[box_num])]
        cv2.rectangle(image,(xmin,ymin),(xmax,ymax),color,5)
        font_scale = 1#min(1,max(3,int(w/500)))
        font_thickness = 2#min(2,max(10,int(w/50)))
        font = cv2.FONT_HERSHEY_SIMPLEX
        # p1,p1 = (int(xmin+1.5*width),ymin), (int(xmin+width),ymin-int(1.5*height))
        text_width = cv2.getTextSize(class_name,font,font_scale,font_thickness)[0][0]
        text_height = cv2.getTextSize(class_name,font,font_scale,font_thickness)[0][1]
        cv2.putText(image,class_name,(xmin,ymin+20),font,font_scale,color,font_thickness)
    names[path] = temp
    return image

```



```

def get_identified_names(image, boxes, labels):
    h,w,_ = image.shape
    temp = []
    for box_num, box in enumerate(boxes):
        xyxy = box.xyxy.int().tolist()
        xmin,ymin,xmax,ymax = xyxy[0][0], xyxy[0][1], xyxy[0][2], xyxy[0][3]
        width = xmax-xmin
        height = ymax-ymin
        class_name = class_names[int(labels[box_num])]
        temp.append(class_name)

    return temp


```

```


COLORS = np.random.uniform(0,255,size=(len(class_names),3))
slices = []
count = 0
d = list()
images_per_row = 5
for j,res in enumerate(results):
    d.append(res)
    if len(d) == images_per_row:
        slices.append(d)
        d = list()

for i in (range(len(slices))):
    fig, axes = plt.subplots(nrows=1, ncols=images_per_row, figsize=(20,10), constrained_layout=True)
    results_ = slices[i]
    for j in range(len(results_)):
        res = results_[j]
        image = res.orig_img
        boxes = res.bboxes
        confidences = boxes.conf
        classes = boxes.cls
        path = res.path
        box_img = plot_box(image, boxes, classes)
        axes[j].imshow(box_img, extent=[0, 30, 0, 30])

```







Part 2 : Tesla Auto pilot and Road Safety

Data wrangling

```
import pandas as pd
import torch
torch.device('cpu')
url = 'https://raw.githubusercontent.com/tksundar/autonomous_driving/refs/heads/master/Tesla%20-%20Deaths.csv'
data = pd.read_csv(url, skip_blank_lines=True, skipinitialspace=True)
data.isna().sum()
data.columns = data.columns.str.rstrip()
nan_counts = data.isna().sum()
columns_to_drop = ['Unnamed: 16', 'Unnamed: 17', 'Source', 'State']
columns_to_drop.extend(nan_counts[nan_counts >= 220].index)
data.drop(columns_to_drop, axis=1, inplace=True)
# Remove rows that does not have a case number
# The values in these rows are outliers
data = data.dropna(subset=['Case #'])
print(data.shape)
print(data.isna().sum())
```

```

(294, 15)
Case #                                0
Year                                0
Date                                0
Country                             0
Description                           0
Deaths                              0
Tesla driver                          5
Tesla occupant                       9
Other vehicle                         4
Cyclists/ Peds                       3
TSLA+cycl / peds                     2
Model                                0
Autopilot claimed                    18
Verified Tesla Autopilot Deaths      4
Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO  1
dtype: int64

```

```

import numpy as np
import matplotlib.pyplot as plt
columns = ['Deaths', 'Tesla driver', 'Tesla occupant', 'Other vehicle', 'Cyclists/ Peds',
           'TSLA+cycl / peds', 'Model', 'Autopilot claimed',
           'Verified Tesla Autopilot Deaths',
           'Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO']
for col in columns:
    data[col] = data[col].apply(lambda x : '0' if x in ['-',' ',''] else x)
    data[col] = pd.to_numeric(data[col], errors='coerce')
    data[col] = data[col].fillna(0)
year_mode = data['Year'].mode()[0]
data['Year'] = data['Year'].apply(lambda x: year_mode if np.isnan(x) else x)
data['Year'] = data['Year'].apply(lambda x: year_mode if x == 202 else x)
data['Year'] = data['Year'].astype(int)
print(data.isna().sum())

```

Case #	0
Year	0
Date	0
Country	0
Description	0
Deaths	0
Tesla driver	0
Tesla occupant	0
Other vehicle	0
Cyclists/ Peds	0
TSLA+cycl / peds	0
Model	0
Autopilot claimed	0
Verified Tesla Autopilot Deaths	0
Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO	0

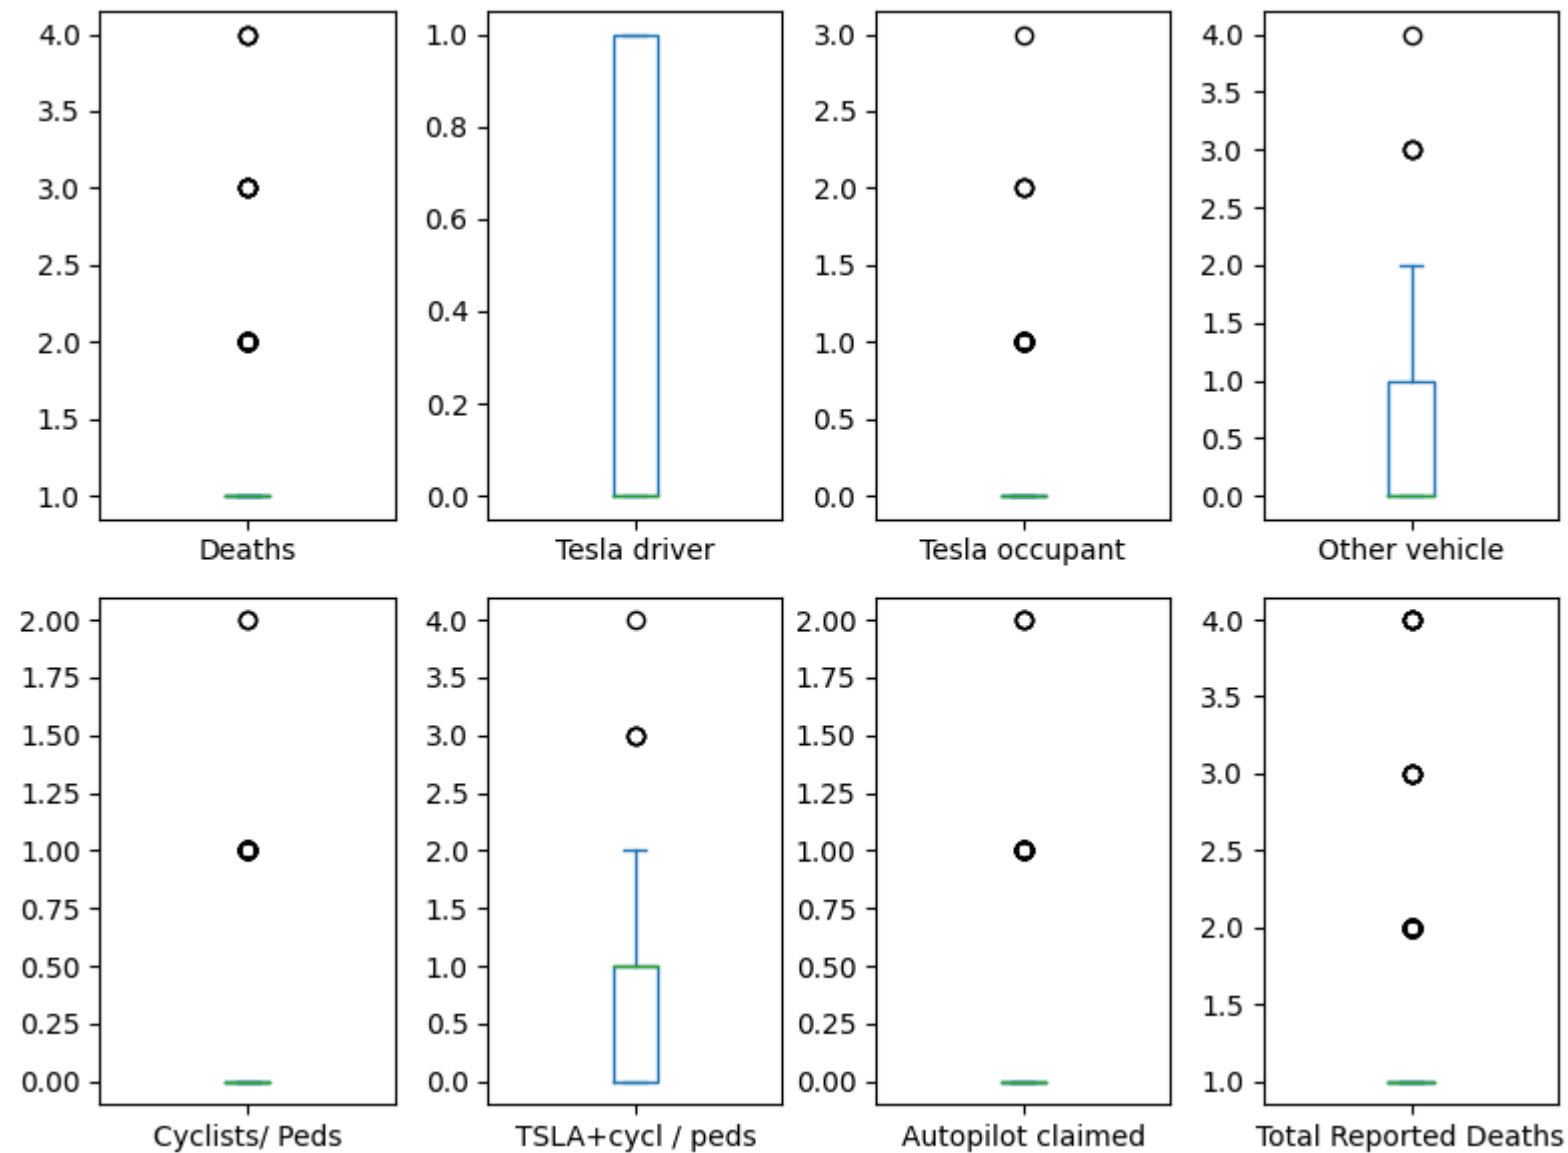
dtype: int64

Distribution of numeric columns

```
data['Country'] = data['Country'].str.rstrip()
data['Date'] = pd.to_datetime(data['Date'])
mod = data['Date'].mode()[0]
data['Date'] = data['Date'].fillna(mod)
data['Total Reported Deaths'] = data.apply(lambda row: row['Deaths']+row['Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA
```

```
columns = ['Deaths', 'Tesla driver', 'Tesla occupant', 'Other vehicle', 'Cyclists/ Peds', 'TSLA+cycl / peds', 'Autopilot claimed', 'Total Reported Deaths']
data[columns].plot(kind='box',
                  subplots=True,
                  layout=(2,4),
                  sharex=False,
                  sharey=False,
                  figsize=(8,6),)

plt.tight_layout()
```




```
data.head()
```


	Case #	Year	Date	Country	Description	Deaths	Tesla driver	Tesla occupant	Other vehicle	Cyclists/ Peds	TSLA+cycl / peds	Model	Autopilot claimed	Verified Tesla Autopilot Deaths	Verified Tesla Autopilot Deaths + All Deaths Reported to NHTSA SGO	To Report Deat
0	294.0	2022	2023-01-17	USA	Tesla crashes into back of semi	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	
1	293.0	2022	2023-01-07	Canada	Tesla crashes	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	
2	292.0	2022	2023-01-07	USA	Tesla hits pole, catches on fire	1.0	0.0	1.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	
3	291.0	2022	2022-12-22	USA	Tesla crashes and burns	1.0	1.0	0.0	0.0	0.0	1.0	0.0	0.0	0.0	0.0	
4	290.0	2022	2022-12-19	Canada	Tesla crashes into storefront	1.0	0.0	0.0	0.0	1.0	1.0	0.0	0.0	0.0	0.0	

Accident Events by country

```
from matplotlib import pyplot as plt
import seaborn as sns
data.groupby('Country').size().plot(kind='barh', color=sns.palettes.mpl_palette('Dark2'))
```

```
plt.gca().spines[['top', 'right']].set_visible(False)
plt.title('Accident Events by country')
```




Deaths by year

```
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')
deaths_by_year = pd.DataFrame(data.groupby('Year')['Total Reported Deaths'].sum()).reset_index()
plt.gca().spines[['top', 'right']].set_visible(False)

fig=sns.barplot(data = deaths_by_year, x='Year', y='Total Reported Deaths' , palette = sns.mpl_palette('Dark2'))
plt.title('Total Deaths by Year')
```

```
plt.yticks(np.arange(20,220,20))
plt.show()
```




Deaths by country and year

```
deaths_by_year = pd.DataFrame(data.groupby(['Country', 'Year'])['Total Reported Deaths'].sum()).reset_index()

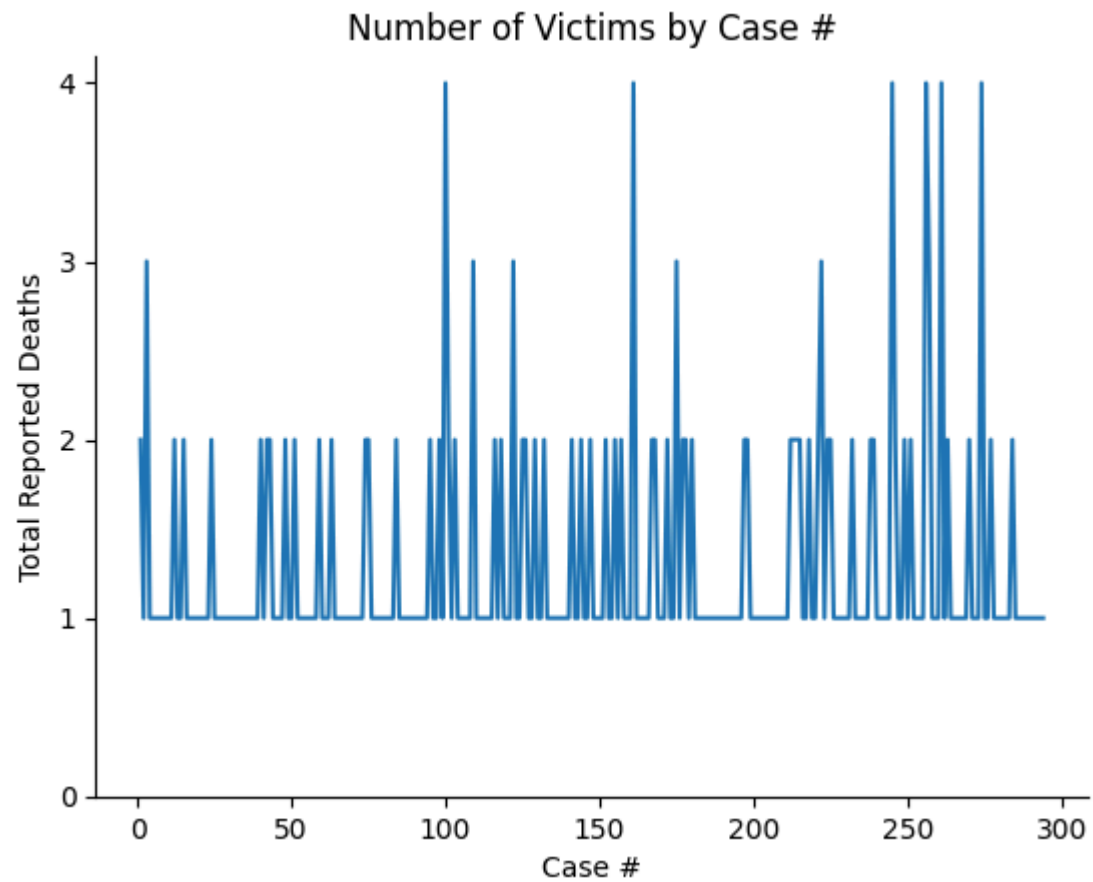
plt.figure(figsize = (12,6))
fig=sns.barplot(data = deaths_by_year, x='Country', hue = "Year", y='Total Reported Deaths', palette = sns.mpl_palette('Set1'))
plt.title('Total Deaths by Country in various years')
plt.xticks(rotation=90)
plt.legend(bbox_to_anchor=(1, 1), loc="upper left")
```

```
plt.gca().spines[['top', 'right']].set_visible(False)
plt.show()
```




Number of victims per accident

```
df = data [['Case #', 'Total Reported Deaths']].sort_values(by='Case #')
sns.lineplot(data=df , x= 'Case #', y='Total Reported Deaths')
plt.gca().spines[['top', 'right']].set_visible(False)
plt.yticks(np.arange(0,5))
_=plt.title('Number of Victims by Case #')
```




Instances of Tesla driver deaths

```
_ = data.groupby('Year')['Tesla driver'].sum().plot(kind='bar', color = sns.mpl_palette('Set1'))  
plt.title('Tesla driver death instances')  
plt.xlabel('Year')  
plt.ylabel('Num Deaths')  
plt.gca().spines[['top', 'right',]].set_visible(False)
```



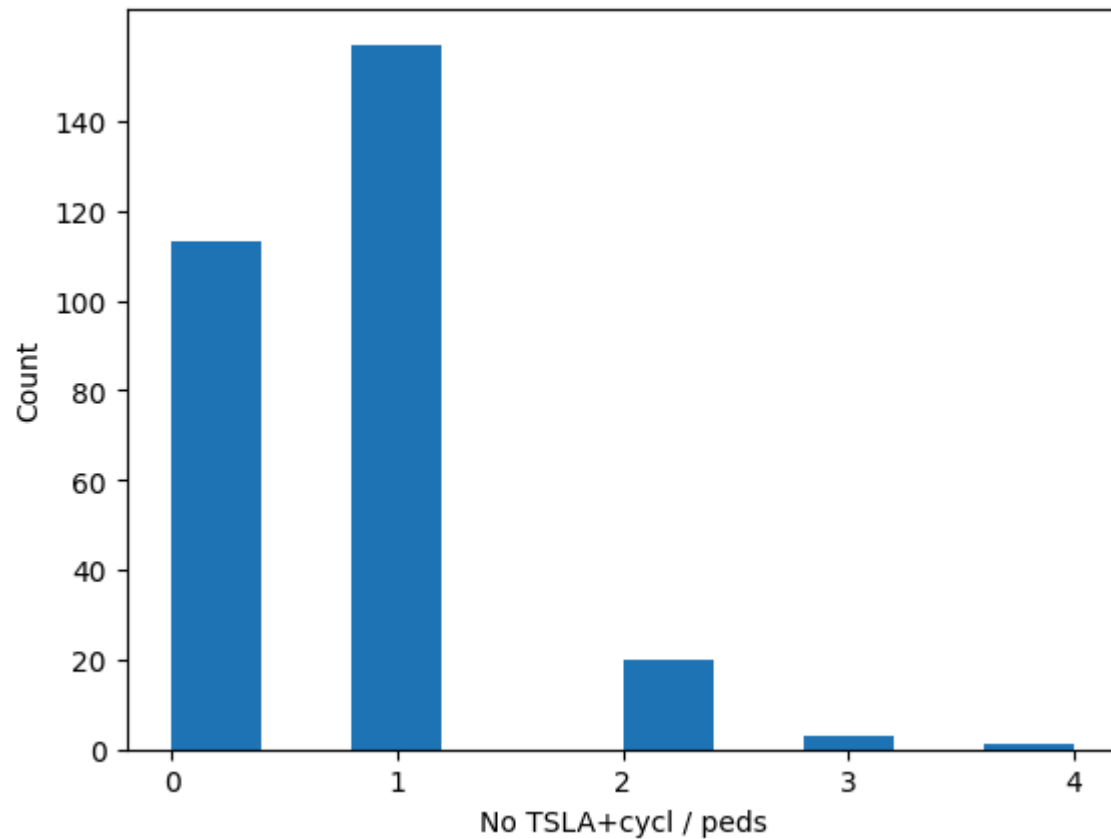
Instances of Tesla occupant deaths

```
data.groupby('Year')['Tesla occupant'].sum().plot(kind='bar', color = sns.mpl_palette('Set1'))  
plt.title('Tesla occupant death instances')  
plt.ylabel('Number of deaths')  
plt.gca().spines[['top', 'right',]].set_visible(False)
```



Distribution of events in which the vehicle hit a cyclist or a pedestrian

```
_ = data['TSLA+cycl / peds'].hist(bins=10, grid=False)
_ = plt.xlabel('No TSLA+cycl / peds')
_ = plt.xticks(np.arange(0, 5, 1))
_ = plt.ylabel('Count')
_ = plt.yticks(np.arange(0, 160, 20))
```




Deaths of occupant, driver of a Tesla along with a cyclist or pedestrian

```
columns = ['Tesla driver', 'Tesla occupant', 'TSLA+cycl / peds']
grp = pd.DataFrame(data.groupby(columns)['Total Reported Deaths'].sum()).reset_index()
grp['Cumulative'] = grp.apply(lambda row: row['Tesla driver'] + row['Tesla occupant'] + row['TSLA+cycl / peds'], axis = 1)
plt.figure(figsize=(10,6))
for i,col in enumerate(columns):
    plt.subplot(1,3,i+1)
    _=sns.barplot(data = grp , x = 'Cumulative',y='Total Reported Deaths',hue=col,palette = sns.mpl_palette('Set1'))
    _=plt.xlabel('Number of deaths')
    plt.gca().spines[['top', 'right']].set_visible(False)
```



```
plt.tight_layout()
```



```
def check_for_two__vehicles(description):
    vehicles = []
    count = 0
    description_lower = description.lower() # For case-insensitive matching
    temp = []
    for object_ in objects:
        if object_.lower() in description_lower:
            temp.append(object_)
```

```
    if(len(temp)>= 2):  
        return 1  
return 0
```

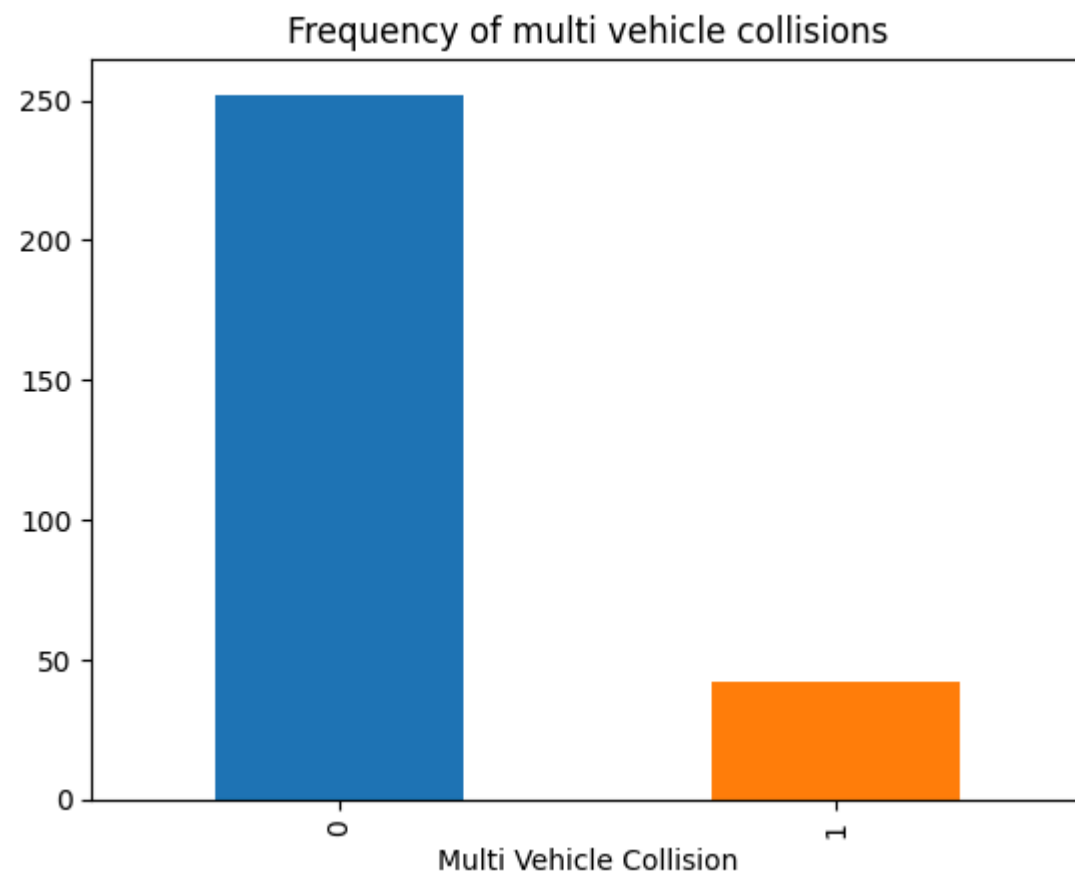
```
url_labels = 'https://raw.githubusercontent.com/tksundar/autonomous_driving/refs/heads/master/labels.csv'  
labels = pd.read_csv(url_labels, names = ['id', 'name', 'xmin', 'ymin', 'xmax', 'ymax'])
```

```
objects = labels['name'].unique().tolist()  
objects.remove('pedestrian')  
objects.append('Tesla')  
objects.append('Semi')
```

Frequency of Tesla colliding with other vehicles

```
data['Multi Vehicle Collision'] = data.apply(lambda row: check_for_two__vehicles(row['Description']),axis=1)  
_=data.groupby(['Multi Vehicle Collision']).size().plot(kind = 'bar', color = sns.mpl_palette('tab10'))  
plt.title('Frequency of multi vehicle collisions')  
  
print(f"Total multi vehicle collisions : {data['Multi Vehicle Collision'].sum()}\n")
```

Total multi vehicle collisions : 42



```
_ = pd.crosstab(data['Other vehicle'], data['Multi Vehicle Collision']).plot(kind='bar')  
plt.ylabel('Number of multi vehicle collisions')  
_ = pd.crosstab(data['Cyclists/ Peds'], data['Multi Vehicle Collision']).plot(kind='bar')  
_ = plt.ylabel('Number of multi vehicle collisions')
```

